



2장

수학적 구성요소

“시도해보지 않고는 누구도 자신이 얼마만큼 해낼 수 있는지 알지 못한다”
푸블릴리우스 시루스



Outline



- ▶ A first example of a neural network
- ▶ Tensors and tensor operations
- ▶ How neural networks learn via backpropagation and gradient descent

○ ○ ○ 2.1 A first look at a neural network ○ ○ ○



Figure 2.1 MNIST
sample digits (28*28)

label: 0 2 4 3

- ▶ classify grayscale images of handwritten digits (28×28 pixels) into their 10 categories (0 through 9)
- ▶ MNIST dataset - a set of 60,000 training images, plus 10,000 test images

Listing 2.1 Loading the MNIST dataset in Keras

```
from keras.datasets import mnist  
  
(train_images, train_labels), (test_images,  
test_labels) = mnist.load_data()
```

○○○ 2.1 *A first look at a neural network* ○○○

- ▶ Let's look at the training data:

```
>>> train_images.shape (60000, 28, 28)
>>> len(train_labels) 60000
>>> train_labels
array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)
```

- ▶ And here's the test data:

```
>>> test_images.shape (10000, 28, 28)
>>> len(test_labels) 10000
>>> test_labels
array([7, 2, 1, ..., 4, 5, 6], dtype=uint8)
```


○○○ 2.1 A first look at a neural network ○○○

1. Network 구성:

Listing 2.2 The network architecture

```
from keras import models
from keras import layers

network = models.Sequential()
network.add(layers.Dense(512, activation='relu',
input_shape=(28 * 28,))) # 784 개 input node
network.add(layers.Dense(10, activation='softmax'))
```



2.1 A first look at a neural network

2. *compilation* step for training:

- A *loss function* —How the network will be able to measure its *performance* on the training data

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \tilde{y}_i)^2$$


$$\text{MAE} = \frac{1}{n} \sum_{j=1}^n |y_j - \hat{y}_j|$$

Cross Entropy Cost Function

$$D(\bar{Y}_i, Y_i) = - \sum Y_i \log \bar{Y}_i$$

$$\begin{bmatrix} \bar{Y}_A \\ \bar{Y}_B \\ \bar{Y}_C \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \begin{bmatrix} Y_A \\ Y_B \\ Y_C \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$$

$$\begin{aligned} - \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \cdot \log \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} &= - \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \begin{bmatrix} 0 \\ \infty \\ \infty \end{bmatrix} \\ &= \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix} \\ &= 0 \end{aligned}$$




2.1 A first look at a neural network

- An *optimizer*—The mechanism through which the network will **update** itself based on the data it sees and its loss function.
- *Metrics to monitor during training and testing*—Here, we'll only care about **accuracy**

Listing 2.3 The compilation step

```
network.compile(optimizer='rmsprop',  
                loss='categorical_crossentropy',  
                metrics=['accuracy'])
```



○○○ 2.1 *A first look at a neural network* ○○○

3. Data Preparation for training:

- scaling - $[0, 255]$ interval $\rightarrow [0, 1]$ interval
- $(60000, 28, 28)$ shape $\rightarrow (60000, 28*28)$ shape

Listing 2.4 Preparing the image data

```
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype('float32') / 255
```


○○○ 2.1 A first look at a neural network ○○○

4. Categorically encode the labels for training:

- One-Hot-Encoding으로 변환

5 \rightarrow [0., 0., 0., 0., 0., 1., 0., 0., 0., 0.], ...

Listing 2.5 Preparing the labels

```
from keras.utils import to_categorical
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

○○○2.1 *A first look at a neural network*○○○

5. **fit** method—we fit the model to its training data :

```
>>> network.fit(train_images, train_labels, epochs=5,
batch_size=128)

Epoch 1/5
60000/60000 [=====] - 9s - loss:
0.2524 - acc: 0.9273
Epoch 2/5
60000/60000 [=====] - ETA: 1s - loss:
0.1035 - acc: 0.9692

...

Epoch 5/5
60000/60000 [=====] - ETA: 12s - loss:
0.0935 - acc: 0.9892
```

○○○ 2.1 *A first look at a neural network* ○○○



6. Test data : a bit lower than the training set accuracy

```
>>> test_loss, test_acc = network.evaluate(test_images, test_labels)
>>> print('test_acc:', test_acc)
test_acc: 0.9785
```

- *overfitting*: the fact that machine-learning models tend to perform worse on new data than on their training data.

2.2 Data representations for neural networks

- ▶ A River Runs Through It - 아름다운 자연을 배경과 함께 강물의 흐름을 따라 고기를 잡는 플라이 낚시와 가족 간의 사랑과 아픔 그리고 삶을 은유
- ▶ 딥러닝과 대단위의 정보와 지식, 이를 통한 추론과 판단이 흐르는 강물과 유사
- ▶ 텐서(tensor) - 0~n 차원(축, axis)까지의 데이터 클래스
 - 0차 텐서: 스칼라(0차원)
 - 1차 텐서: 벡터(1차원)
 - 2차 텐서: 행렬(2차원),
 - 3차원 이상: n차 텐서
- ▶ 텐서가 입력에서 출력까지 흐르며 학습



2.2 Data representations for neural networks

1. Scalars (0D tensors)

- ▶ A **scalar tensor** contains only **one number**
- ▶ In Numpy, a **float32** or **float64** number is a **scalar tensor** (or **scalar array**).
- ▶ **ndim** attribute - the number of axes in Numpy tensor
- ▶ a **scalar tensor** - 0 **axes** (or **rank**), (`ndim == 0`)
- ▶ Here's a Numpy scalar:

```
>>> import numpy as np
>>> x = np.array(12)    # scalar tensor
>>> x
array(12)
>>> x.ndim
0
```

2.2 Data representations for neural networks

2. Vectors (1D tensors)

- ▶ An array of numbers is called a **vector**, or **1D tensor**.
- ▶ A **1D tensor** is said to have exactly one **axis**.
- ▶ Following is a Numpy **vector**:

```
>>> x = np.array([12, 3, 6, 14])
```

```
>>> x
```

```
array([12, 3, 6, 14])
```

```
>>> x.ndim
```

```
1
```

2.2 Data representations for neural networks

3. Matrices (2D tensors)

- ▶ An array of vectors is called *matrix*, or **2D tensor**.
- ▶ A **matrix** has two axes (often referred to *rows* and *columns*).
- ▶ This is a Numpy matrix:

```
>>> x = np.array([[5, 78, 2, 34, 0], # first row of x
                  [6, 79, 3, 35, 1],
                  [7, 80, 4, 36, 2]])

>>> x.ndim # first column of x
2
```

2.2 Data representations for neural networks

4. 3D tensors and higher-dimensional tensors (nD tensors)

► If you pack such matrices in a new array, you obtain a 3D tensor, which you can visually interpret as a cube of numbers. Following is a Numpy 3D tensor:

```
>>> x = np.array([ [ [5, 78, 2, 34, 0],  
                    [6, 79, 3, 35, 1],  
                    [7, 80, 4, 36, 2] ],  
                  [ [5, 78, 2, 34, 0],  
                    [6, 79, 3, 35, 1],  
                    [7, 80, 4, 36, 2] ],  
                  [ [5, 78, 2, 34, 0],  
                    [6, 79, 3, 35, 1],  
                    [7, 80, 4, 36, 2] ] ] )
```

```
>>> x.ndim  
3
```


2.2 Data representations for neural networks

5. Key attributes – tensor의 정의

- Number of *axes* (*rank*) — 3D tensor == 3 axes, matrix == 2 axes, `ndim` in Numpy
- *Shape* — a tuple of integers
 - shape of a scalar - `()`
 - shape of a vector - `(5,)`
 - shape of 3D tensor - `(3, 3, 5)`
- *Data type* (usually called `dtype` in Python libraries) — type of the data contained in the tensor;
 - ex. `float32`, `uint8`, `float64`, `char`, no string tensors

2.2 Data representations for neural networks

5. Key attributes – tensor의 정의

- ▶ The data in the MNIST dataset:

```
from keras.datasets import mnist  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

- ▶ The number of axes of the tensor `train_images`, the `ndim` attribute:

```
>>> print(train_images.ndim)  
3
```

- ▶ Here's its shape:

```
>>> print(train_images.shape)  
(60000, 28, 28)
```

- ▶ And this is its data type, the `dtype` attribute:

```
>>> print(train_images.dtype)  
uint8
```

2.2 Data representations for neural networks

5. Key attributes – tensor의 정의

▶ Let's display the fourth digit in this 3D tensor, using the library Matplotlib; see figure 2.2.

Listing 2.6 Displaying the fourth digit

```
digit = train_images[4]
import matplotlib.pyplot as plt
plt.imshow(digit, cmap=plt.cm.binary)
plt.show()
```




Figure 2.2 The fourth sample in our dataset

2.2 Data representations for neural networks

6. Manipulating tensors in Numpy

- ▶ *tensor slicing* - Selecting specific elements in a tensor

```
train_images[i]
```

- ▶ selects digits #10 to #100 (#100 isn't included, 10부터 100개) :

```
>>> my_slice = train_images[10:100]
```

```
>>> print(my_slice.shape) (90, 28, 28)
```

```
>>> my_slice = train_images[10:100, :, :]
```

```
>>> my_slice.shape
```

```
(90, 28, 28)
```

```
>>> my_slice = train_images[10:100, 0:28, 0:28]
```


```
>>> my_slice.shape
```

```
(90, 28, 28)
```

14 × 14 pixels centered in the middle

```
>>> my_slice = train_images[:, 7:-7, 7:-7]
```

7번째부터 끝에서 8번째까지



2.2 Data representations for neural networks

7. The notion of data batches

- ▶ here's one batch of our MNIST digits, with batch size of 128:
- ▶ the **first axis (axis 0)** is called the *batch axis* or *batch dimension*

```
batch = train_images[:128] # 0-127  
batch = train_images[128:256] # 127-255
```

- the n th batch:

```
batch = train_images[128*n : 128*(n + 1)]
```

2.2 Data representations for neural networks

8. Real-world examples of data tensors

- *Vector data*—2D tensors of shape (samples, features)
- *Timeseries data or sequence data*—3D tensors of shape (samples, timesteps, features)
- *Images*—4D tensors of shape (samples, height, width, channels) or (samples, channels, height, width)
- *Video*—5D tensors of shape (samples, frames, height, width, channels) or (samples, frames, channels, height, width)

2.2 Data representations for neural networks

9. Vector data

the first axis is the *samples axis* and the second axis is the *features axis*

- dataset of people - age, ZIP code, and income.
100,000 people - 2D tensor of shape (100000, 3)
- dataset of text documents - each document by the counts of how many times each word appears in it (out of a dictionary of 20,000 common words)
500 documents - 2D tensor of shape (500, 20000).

2.2 Data representations for neural networks

10. Timeseries data or sequence data

- A dataset of stock prices - 250 days , 390 minutes in a trading day, 3 features in a 3D tensor of shape $(250, 390, 3)$




Figure 2.3 A 3D timeseries data tensor

- A dataset of tweets - 280 characters out of an alphabet of 128 unique characters $(280, 128)$,
dataset of 1 million tweets - $(1000000, 280, 128)$

2.2 Data representations for neural networks

11. Image data

- a batch of 128 color images could be stored in a tensor of shape $(128, 256, 256, 3)$ (see figure 2.4)
- $(\text{samples}, \text{height}, \text{width}, \text{color_depth})$.




Figure 2.4 A 4D image data tensor (channels-first convention)

2.2 Data representations for neural networks

12. Video data

- a batch of different videos can be stored in a 5D tensor of shape (samples, frames, height, width, color_depth)
- a 60-second, 144×256 YouTube video clip sampled at 4 frames per second would have 240 frames –
(4, 240, 144, 256, 3)

2.3 The gears of neural networks: tensor operations

- ▶ building our network by stacking Dense layers

- ▶ A Keras layer instance looks like this:

```
keras.layers.Dense(512,  
activation='relu')
```

- ▶ where W is a 2D tensor and b is a vector, both attributes of the layer:

$$\text{output} = \text{relu}(\text{dot}(W, \text{input}) + b)$$

2.3 The gears of neural networks: tensor operations

1. Element-wise operations

Matrix vs. Element-wise operations

```
>> A=[1 2; 4 5]      >> B = [10 20; 30 40]
```

1	2	10	20
4	5	30	40

- Matrix multiplication

```
>> A * B
```

70	100
190	280

- Element-wise multiplication

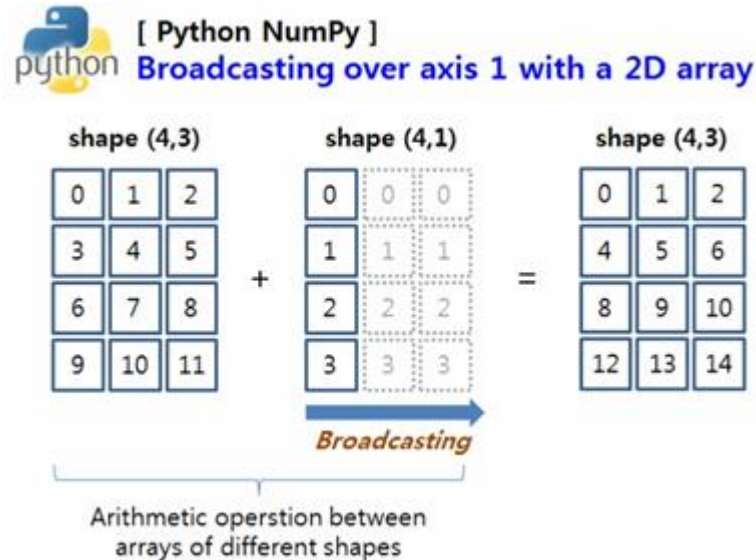
```
>> A .* B
```

10	40
120	200

```
import numpy as np  
z = x + y  
z = np.maximum(z, 0.)
```

2.3 The gears of neural networks: tensor operations

2. Broadcasting



- ▶ element-wise **maximum** operation to two tensors of different shapes via broadcasting:

```
import numpy as np
x = np.random.random((64, 3, 32, 10))
y = np.random.random((32, 10))
z = np.maximum(x, y) # shape (64, 3, 32, 10) like x.
```

2.3 The gears of neural networks: tensor operations

3. Tensor dot

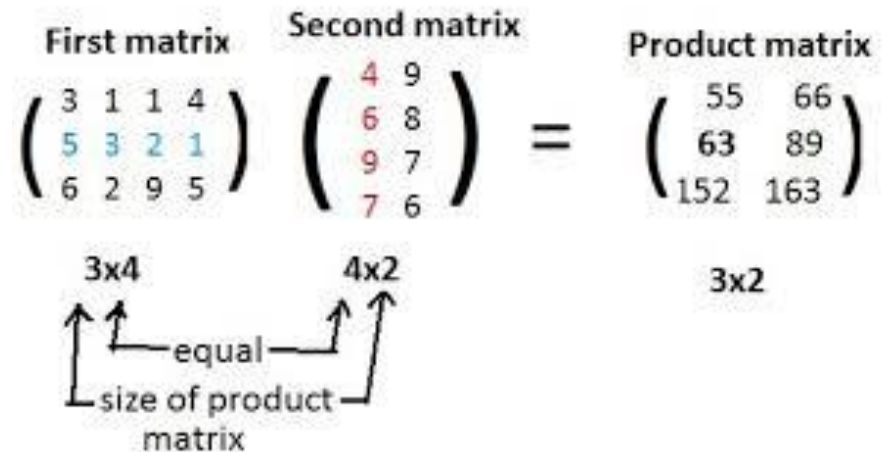



Figure 2.5 Matrix dot-product box diagram

2.3 The gears of neural networks: tensor operations

4. Tensor reshaping

```
train_images = train_images.reshape((60000, 28 * 28))
>>> x = np.array([[0., 1.],
                  [2., 3.],
                  [4., 5.]])

>>> print(x.shape)
(3, 2)
>>> x = x.reshape((6, 1))
>>> x
array([[ 0.],
       [ 1.],
       [ 2.],
       [ 3.],
       [ 4.],
       [ 5.]])

>>> x = x.reshape((2, 3))
>>> x
array([[ 0.,  1.,  2.],
       [ 3.,  4.,  5.]])
```

► A special case of reshaping that's commonly encountered is *transposition*. *Transposing* a matrix means exchanging its rows and its columns, so that $x[i, :]$ becomes $x[:, i]$:

```
>>> x = np.zeros((300, 20))
>>> x = np.transpose(x)
>>> print(x.shape)
(20, 300)
```

2.4 The engine of neural networks: gradient-based optimization

1. What's a derivative?

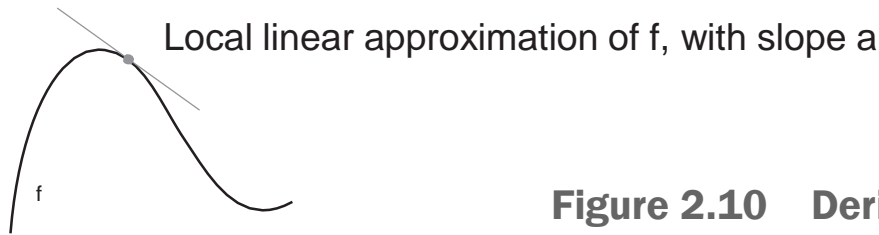



Figure 2.10 Derivative of f in p

- **Derivative:** A *gradient* is the derivative of a tensor operation.

2.4 The engine of neural networks: gradient-based optimization

3. Stochastic gradient descent

- ▶ The term *stochastic* refers to the fact that each batch of data is drawn at random.
- 1 Draw a **batch** of training samples x and corresponding targets y .
- 2 Run the network on x to obtain **predictions** y_{pred} .
- 3 Compute the **loss** of the network on the batch, a measure of the mismatch between y_{pred} and y .
- 4 Compute the **gradient of the loss** with regard to the network's parameters (a *backward pass*).
- 5 Move the parameters a little in the **opposite direction** from the gradient—for example $W \leftarrow W - \text{step} * \text{gradient}$ —thus **reducing the loss** on the batch a bit.



- ▶ If step is too small, the descent down the curve will take many iterations - stuck in a local minimum.
- ▶ If step is too large - completely random locations on the curve.

Figure 2.11 SGD down a 1D loss curve (one learnable parameter)

2.4 The engine of neural networks: gradient-based optimization

3. Stochastic gradient descent

- ▶ visualize gradient descent along a 2D loss surface, as shown in figure 2.12.
- ▶ can't possibly visualize what the actual process of training a neural network looks like—1,000,000-dimensional space
- ▶ *Optimizers* - Adagrad, RMSProp, ...




Figure 2.12 Gradient descent down a 2D loss surface (two learnable parameters)

2.4 The engine of neural networks: gradient-based optimization

3. Stochastic gradient descent

- **Momentum** with SGD: convergence speed and local minima




Figure 2.13 A local minimum and a global minimum

2.4 The engine of neural networks: gradient-based optimization

5. Looking back at our first example

- ▶ review each piece of it in the light of what you've learned in the previous three sections

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype('float32') / 255
```

```
network = models.Sequential()
network.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
network.add(layers.Dense(10, activation='softmax'))
```

```
network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
network.fit(train_images, train_labels, epochs=5, batch_size=128)
```