



# 3장 *Getting started with neural networks*

“기회와 준비가 만났을 때 ... ”



## 5. Classifying newswires: a multiclass classification example



- ▶ more than two classes?
- ▶ classify Reuters newswires into 46 mutually exclusive topics - *multi-class classification*

### 3.5.1 The Reuters dataset

- ▶ a set of short newswires and their topics, published by Reuters in 1986 for text classification
- ▶ 46 different topics; each topic has at least 10 examples in the training set
- ▶ Like IMDB and MNIST, the Reuters dataset comes packaged as part of Keras.

## 5. Classifying newswires: a multiclass classification example

### Listing 3.12 Loading the Reuters

```
from keras.datasets import reuters
(train_data, train_labels), (test_data, test_labels) =
    reuters.load_data(num_words=10000)
```

▶ `num_words=10000` restricts the data to the 10,000 most frequently occurring words found in the data

▶ 8,982 training examples and 2,246 test examples:

```
>>> len(train_data) 8982
```

```
>>> len(test_data) 2246
```

▶ each example is a list of integers (word indices):

```
>>> train_data[10]
[1, 245, 273, 207, 156, 53, 74, 160, 26, 14, 46, 296, 26, 39, 74, 2979, 3554,
14, 46, 4689, 4329, 86, 61, 3499, 4795, 14, 61, 451, 4329, 17, 12]
```

▶ `label : 0 and 45`, a topic index

```
>>> train_labels[10]
```

```
3
```



## 5. Classifying newswires: a multiclass classification example



### 3.5.2 Preparing the data

► vectorize the data - 46 different topics; each topic has at least 10 examples in the training set

#### Listing 3.14 Encoding the data

```
import numpy as np
```

```
def vectorize_sequences(sequences, dimension=10000):  
    results = np.zeros((len(sequences), dimension))  
    for i, sequence in enumerate(sequences):  
        results[i, sequence] = 1.  
    return results
```

```
x_train = vectorize_sequences(train_data)
```

```
x_test = vectorize_sequences(test_data)
```

근데 책에는 원핫 인코딩인데 교수님은 Binary Metrics라고 말하심  
원핫 인코딩을 통해 10000차원 벡터로 변환



## 5. Classifying newswires: a multiclass classification example



### 3.5.2 Preparing the data

► **One-hot encoding** (*categorical encoding*) - an all-zero vector with a 1 in the place of the label index. Here's an example:

```
def to_one_hot(labels, dimension=46):  
    results = np.zeros((len(labels), dimension))  
    for i, label in enumerate(labels):  
        results[i, label] = 1.  
    return results  
one_hot_train_labels = to_one_hot(train_labels)  
one_hot_test_labels = to_one_hot(test_labels)
```

► built-in way to do this in Keras, in the MNIST example:

```
from keras.utils.np_utils import to_categorical  
  
one_hot_train_labels = to_categorical(train_labels)  
one_hot_test_labels = to_categorical(test_labels)
```

## 5. Classifying newswires: a multiclass classification example

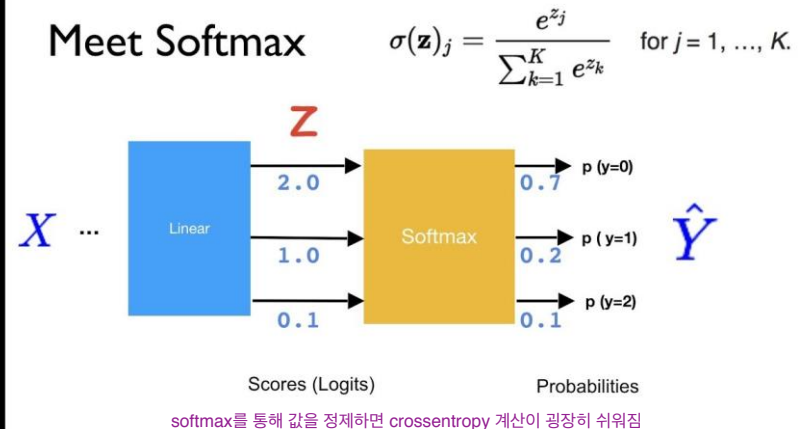
### 3.5.3 Building your network

- ▶ the number of output classes - 46
- ▶ use larger layers with 64 units

#### Listing 3.15 Model definition

```
from keras import models
from keras import layers
model = models.Sequential()
model.add(layers.Dense(64, activation='relu',
                        input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))
```

- ▶ There are two other things you should note about this architecture:
  - the network will output a 46-dimensional vector
  - The last layer uses a **softmax** activation - a *probability distribution* over the 46 different output classes — `output[i]` is the probability that the sample belongs to class `i`. The 46 scores will sum to 1.



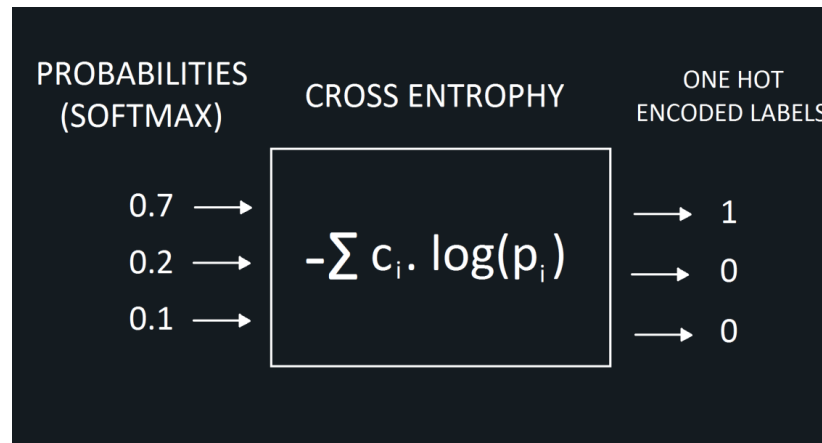
## 5. Classifying newswires: a multiclass classification example

### 3.5.3 Building your network

- ▶ loss function - `categorical_crossentropy`, distance between the probability distribution **output** by the network and the true distribution of the **labels**, minimizing the distance

#### Listing 3.16 Compiling the model

```
model.compile(optimizer='rmsprop',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])
```





## 5. Classifying newswires: a multiclass classification example



### 3.5.4 Validating your approach

validation은 왜 만드나?  
- 이것도해보고 저것도해보고싶어서  
- train->test하면 시간이 많이 지체되니까 모의로 하는 것

- ▶ Let's set apart 1,000 samples in the training data to use as a validation set.

#### Listing 3.17 Setting aside a validation set

```
x_val = x_train[:1000] partial_x_train = x_train[1000:]  
y_val = one_hot_train_labels[:1000] partial_y_train =  
        one_hot_train_labels[1000:]
```

validation에 1000개      나머지 7982개는 train

- ▶ Now, let's train the network for 20 epochs.

#### Listing 3.18 Training the model

```
history = model.fit(partial_x_train,  
                    partial_y_train, epochs=20,  
                    batch_size=512,  
                    validation_data=(x_val, y_val))
```



## 5. Classifying newswires: a multiclass classification example

### 3.5.4 Validating your approach

► let's display its **loss** and **accuracy** curves (see figures 3.9 and 3.10)

#### Listing 3.19 Plotting the training and validation loss

```
import matplotlib.pyplot as plt
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(loss) + 1)
plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




Figure 3.9 Training and validation **loss**

## 5. Classifying newswires: a multiclass classification example

### 3.5.4 Validating your approach

► let's display its **loss** and **accuracy** curves (see figures 3.9 and 3.10)

#### **Listing 3.20** Plotting the training and validation accuracy

```
plt.clf()
acc = history.history['acc']
val_acc = history.history['val_acc']
plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```




Figure 3.10 Training and validation **accuracy**

## 5. Classifying newswires: a multiclass classification example

### 3.5.4 Validating your approach

이거 오버피팅이 일어나는지 어떻게 구분하나?

▶ The network begins to **overfit after nine epochs**. Let's train a new network from scratch for nine epochs and then evaluate it on the test set.

#### **Listing 3.21 Retraining a model from scratch**

```
model = models.Sequential()
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(46, activation='softmax'))

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(partial_x_train,
          partial_y_train, epochs=9, batch_size=512,
          validation_data=(x_val, y_val))
results = model.evaluate(x_test, one_hot_test_labels)
```

▶ Here are the final results:

```
>>> results
[0.9565213431445807, 0.79697239536954589]
```



## 5. Classifying newswires: a multiclass classification example



### 3.5.5 Generating predictions on new data

▶ You can verify that the `predict` method of the model instance returns a probability distribution over all 46 topics. Let's generate topic predictions for all of the test data.

#### **Listing 3.22** Generating predictions for new data

```
predictions = model.predict(x_test)
```

▶ Each entry in `predictions` is a vector of length 46:

```
>>> predictions[0].shape  
(46,)
```

▶ The coefficients in this vector **sum** to 1:

```
>>> np.sum(predictions[0])  
1.0
```

The **largest entry** is the **predicted class**:

```
>>> np.argmax(predictions[0]) # 1.00이 있는 위치의 index 값  
4
```



## 5. Classifying newswires: a multiclass classification example



### 3.5.6 A different way to handle the labels and the loss

- ▶ We mentioned earlier that another way to encode the labels would be to cast them as an **integer** tensor, like this:

```
y_train = np.array(train_labels)
y_test = np.array(test_labels)
```

- ▶ loss function –

```
categorical_crossentropy → sparse를 쓰면 one-hot 인코딩을 안해도 알아서 해준다? sparse_categorical_crossentropy:
model.compile(optimizer='rmsprop',
               loss='sparse_categorical_crossentropy',
               metrics=['acc'])
```

## 5. Classifying newswires: a multiclass classification example

### 3.5.7 The importance of having sufficiently large intermediate layers

▶ final outputs are 46-dimensional, intermediate layers << 46-dimensional: for example, 4-dimensional.

#### Listing 3.23 A model with an information bottleneck

```
model = models.Sequential() input이 크게 들어왔는데 4로 좁혀졌다가 다시 커짐 -> 병목 현상
model.add(layers.Dense(64, activation='relu', input_shape=(10000,)))
model.add(layers.Dense(4, activation='relu')) # too low-dimensional → 8 ...
model.add(layers.Dense(46, activation='softmax')) # hyperplanes of 46 classes
model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(partial_x_train,
          partial_y_train, epochs=20, batch_size=128,
          validation_data=(x_val, y_val))
```

▶ The network now peaks at ~71% validation accuracy, an 8% absolute drop.

## 5. Classifying newswires: a multiclass classification example

### 3.5.8 Further experiments

- Try using larger or smaller layers: **32 units**, **128 units**, and so on.
- You used two hidden layers. Now try using a **single** hidden layer, or **three** hidden layers.

### 3.5.9 Wrapping up

- ▶ Here's what you should take away from this example:
  - **classify  $N$  classes**, Dense **output layer** of size  $N$ .
  - single-label (select one), multiclass classification problem - **softmax** activation to output a probability distribution over the  $N$  output classes.
  - **Categorical crossentropy** - minimizes the distance between the probability distributions output
  - There are two ways to handle labels in multiclass classification:
    - **one-hot encoding** using `categorical_crossentropy` as a loss function
    - **the labels as integers** using the `sparse_categorical_crossentropy` loss function
  - multiclass classification problem - avoid creating information **bottlenecks**