



3장 *Getting started with neural networks*

“기회와 준비가 만났을 때 ...”

6. Predicting house prices: a regression example



<-> Classification

회귀 모델

- ▶ The two previous examples - predict a **single discrete label** of an input data point.
- ▶ Another common type of machine-learning problem is **regression** - predicting a **continuous value**: temperature tomorrow, time, or price

NOTE Don't confuse *regression* and the algorithm *logistic regression*. Confusingly, logistic regression isn't a regression algorithm—it's a classification algorithm.

regression 이라고해서 다 regression은 아니고
Linear나 Logistic이 붙으면 Classification이다!



6. Predicting house prices: a regression example

3.6.1 The Boston Housing Price dataset

- ▶ predict the median price of homes in a given Boston suburb with the crime rate in the mid-1970s
- ▶ only 506 samples - 404 training samples and 102 test samples.
- ▶ And each *feature* in the input data (for example, the crime rate, the local property tax rate) has a **different scale**. For instance, some values are proportions, which take values between **0 and 1**; others take values between **1 and 12**, others between **0 and 100**, and so on.

▶ **Listing 3.24 Loading the Boston housing**

```
from keras.datasets import boston_housing
(train_data, train_targets), (test_data, test_targets)
    =boston_housing.load_data()
```

- ▶ Let's look at the data:

```
>>> train_data.shape
(404, 13) # 13 features
>>> test_data.shape
(102, 13)
```



6. Predicting house prices: a regression example



3.6.1 The Boston Housing Price dataset

▶ As you can see, you have 404 training samples and 102 test samples, each with **13 numerical features** - capita crime rate, average number of rooms per dwelling, accessibility to highways, and so on.

▶ The targets are the median values of owner-occupied homes, in thousands of dollars:

```
>>> train_targets  
[15.2, 42.3, 50. ... 19.4, 19.4, 29.1] # 404 in $(*1000)
```

▶ The prices are typically between \$10,000 and \$50,000

6. Predicting house prices: a regression example



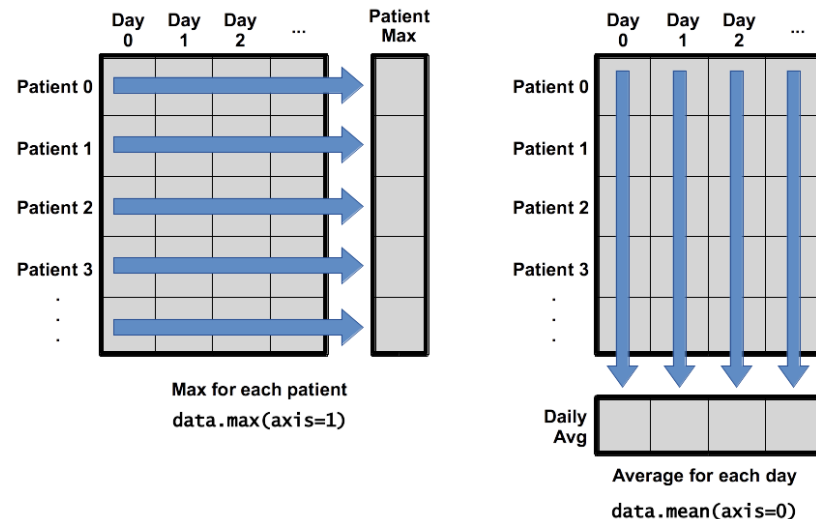
3.6.2 Preparing the data

- ▶ **feature-wise normalization**: feature is centered around 0 and has a unit standard deviation.

Listing 3.25 Normalizing the data

```
mean = train_data.mean(axis=0) # (404,13)
train_data -= mean
std = train_data.std(axis=0) # (102,13)
train_data /= std
test_data -= mean
test_data /= std
```

- ▶ Note that the quantities used for normalizing the test data are computed using the training data.





6. Predicting house prices: a regression example



3.6.3 Building your network

- ▶ small samples - small network with two hidden layers, each with 64 units.
- ▶ **less training data** - **worse overfitting**, a **small network** is one way to mitigate overfitting.

Listing 3.26 Model definition

```
from keras import models
from keras import layers
def build_model():
    model = models.Sequential()
    model.add(layers.Dense(64, activation='relu',
        input_shape=(train_data.shape[1],)))
    model.add(layers.Dense(64,
        activation='relu'))
    model.add(layers.Dense(1))
    model.compile(optimizer='rmsprop', loss='mse', metrics=['mae'])
    return model
```

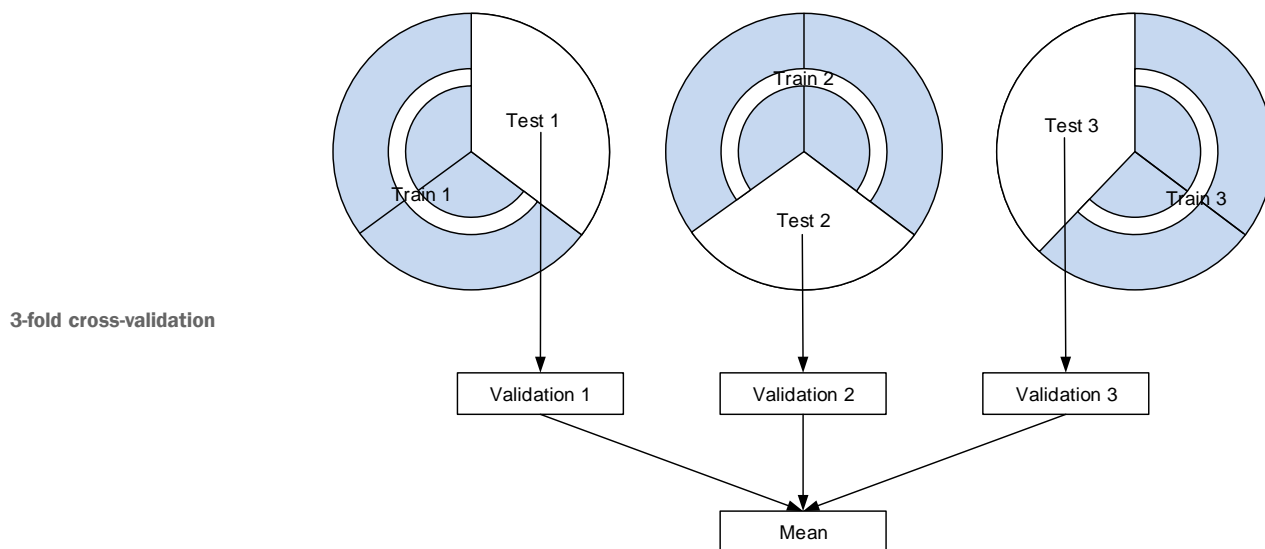
- ▶ output - a **single unit** and **no activation** (it will be a linear layer) for scalar **regression** to predict a **single continuous value**, free to learn to predict values in any range
- ▶ sigmoid activation function - predict values between 0 and 1
- ▶ *mean absolute error* (**MAE**) loss function for regression problems - 0.5 → off by \$500 on average.

6. Predicting house prices: a regression example

3.6.4 Validating your approach using *K*-fold validation

매우 중요(Validation하는 방법)

- ▶ very **small validation set** (for instance, about 100 examples) - **high variance** with regard to the validation split
- ▶ ***K*-fold cross-validation** (see figure) - splitting the available data into K partitions (typically $K = 4$ or 5), instantiating K identical models, and training each one on $K - 1$ partitions while evaluating on the remaining partition.
- ▶ The **validation score** for the model - average of the K validation scores





6. Predicting house prices: a regression example



3.6.4 Validating your approach using K-fold validation

Listing 3.27 K-fold validation

```
import numpy as np
k = 4
num_val_samples = len(train_data) // k  # 나눗셈의 몫
num_epochs = 100
all_scores = []
for i in range(k):  # i = 0, 1, 2, 3
    print('processing fold #', i)
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]
    partial_train_data = np.concatenate([train_data[:i * num_val_samples],
                                          train_data[(i + 1) * num_val_samples:]],
                                          axis=0)
    partial_train_targets = np.concatenate([train_targets[:i * num_val_samples],
                                           train_targets[(i + 1) * num_val_samples:]],
                                           axis=0)
    model = build_model()
    history = model.fit(partial_train_data, partial_train_targets,
                        validation_data=(val_data, val_targets), epochs=num_epochs,
                        batch_size=1, verbose=0)
    mae_history = history.history['val_mean_absolute_error']
    all_mae_histories.append(mae_history)
```




6. Predicting house prices: a regression example



3.6.4 Validating your approach using K-fold validation

- ▶ Running this with `num_epochs = 100` yields the following results:

```
>>> all_scores  
[2.588258957792037, 3.1289568449719116, 3.1856116051248984, 3.0763342615401386]  
>>> np.mean(all_scores)  
2.9947904173572462
```

- ▶ different validation scores, from 2.6 to 3.2.
- ▶ The average (3.0) is a much more reliable metric **K-fold cross-validation**
- ▶ \$3,000 on average - significant considering with \$10,000 to \$50,000.

6. Predicting house prices: a regression example

3.6.4 Validating your approach using K-fold validation

► Let's try training the network a bit longer: **500 epochs**. To keep a record of how well the model does at each epoch, you'll modify the training loop to save the per-epoch validation score log.

Listing 3.28 Saving the validation logs

```
num_epochs = 500
all_mae_histories = []
for i in range(k):
    print('처리중인 폴드 #', i)
    val_data = train_data[i * num_val_samples: (i + 1) * num_val_samples]
    val_targets = train_targets[i * num_val_samples: (i + 1) * num_val_samples]
    partial_train_data = np.concatenate(
        [train_data[:i * num_val_samples],
         train_data[(i + 1) * num_val_samples:]],
        axis=0)
    partial_train_targets = np.concatenate(
        [train_targets[:i * num_val_samples],
         train_targets[(i + 1) * num_val_samples:]],
        axis=0)
    model = build_model()
    history = model.fit(partial_train_data, partial_train_targets,
                        validation_data=(val_data, val_targets),
                        epochs=num_epochs, batch_size=1, verbose=0)
    mae_history = history.history['val_mean_absolute_error']
    all_mae_histories.append(mae_history)
```



6. Predicting house prices: a regression example



3.6.4 Validating your approach using K-fold validation

- ▶ the average of the per-epoch MAE scores for all folds.


Listing 3.29 Building the history of successive

```
average_mae_history = [np.mean([x[i] for x in all_mae_histories])  
                        for i in range(num_epochs)]
```

Let's plot this; see figure 3.12.

Listing 3.30 Plotting validation scores

```
import matplotlib.pyplot as plt  
plt.plot(range(1, len(average_mae_history) + 1),  
         average_mae_history)  
plt.xlabel('Epochs')  
plt.ylabel('Validation MAE')  
plt.show()
```



이 다음에 그래프 나오는 ppt 한장이 찢림

Figure 3.12 Validation MAE by epoch



6. Predicting house prices: a regression example



- ▶ validation MAE stops improving significantly after 80 epochs.
- ▶ adjust the size of the hidden layers, and then look at its performance on the test data.

Listing 3.32 Training the final model

```
model = build_model()  
model.fit(train_data, train_targets, epochs=80, batch_size=16, verbose=0)  
test_mse_score, test_mae_score = model.evaluate(test_data, test_targets)
```

- ▶ Here's the final result:

```
>>> test_mae_score 2.5532484335057877
```

- ▶ You're still off by about \$2,550



6. Predicting house prices: a regression example



3.6.5 Wrapping up

- ▶ Here's what you should take away from this example:
 - Mean squared error (**MSE**) is a **loss function** commonly used for **regression**.
 - The concept of **accuracy** doesn't apply for regression. A common regression **metric** is mean absolute error (**MAE**). output 노드의 개수는 1개, activation 함수 따로 없음
 - When **features** in the input data have values in **different ranges**, each feature should be **scaled** independently as a preprocessing step.
 - When there is **little data** available, using **K-fold validation** is a great way to reliably evaluate a model.
 - When **little training data** is available, it's preferable to use a small network with **few hidden layers** (typically only **1 or 2**), in order to avoid severe **overfitting**.



6. Predicting house prices: a regression example



Chapter summary

- binary classification, multiclass classification, and scalar regression
- preprocess raw data before feeding it into a neural network.
- features with different ranges, scale each feature independently
- As training progresses, neural networks eventually begin to overfit on never-before-seen data.
- If you have small training data, use a small network with only one or two hidden layers, to avoid severe overfitting.
- If your data is divided into many categories, you may cause information bottlenecks if you make the intermediate layers too small.
- Regression uses different loss functions and different evaluation metrics than classification.
- When you're working with little data, K-fold validation can help reliably evaluate your model.