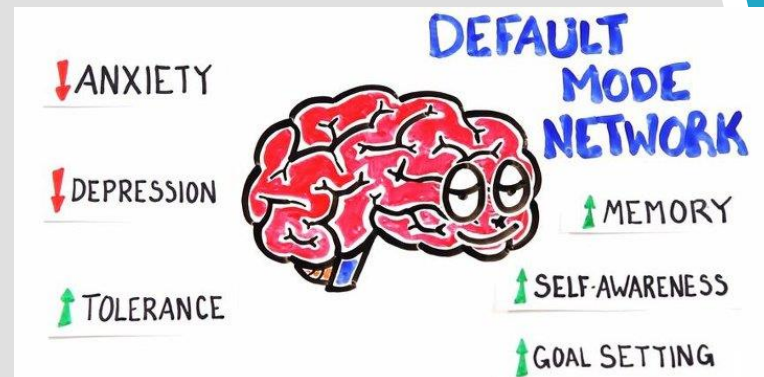



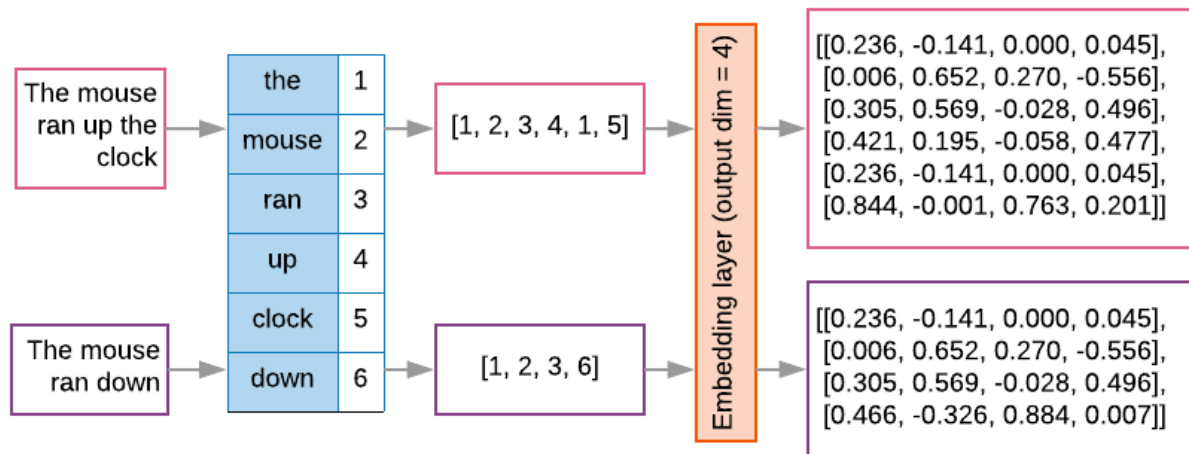
6장 *Deep learning for text and sequences*

“default mode network”




6.2 Understanding recurrent neural networks

- ▶ A **major characteristic** of all neural networks is that they have **no memory** - *no state kept in between inputs*
- ▶ **feedforward networks** - IMDB example: an entire movie review was transformed into a **single large vector** and processed in one go.



6.2 Understanding recurrent neural networks

- ▶ **word by word** (eye saccade by eye saccade) - from **past information** → constantly updated **as new information**
- ▶ A **recurrent neural network** (RNN) - it processes sequences by **iterating** through the sequence elements and maintaining a **state** containing information relative to what it has seen so far.



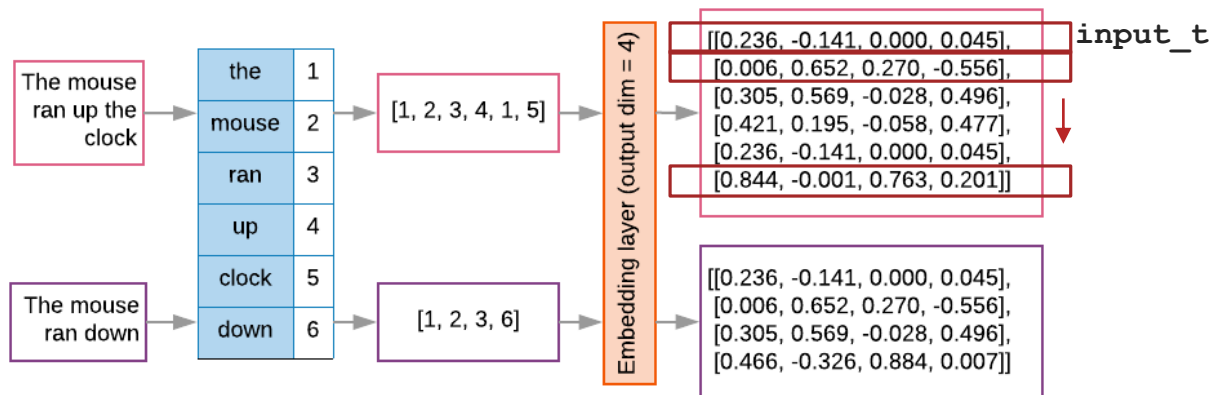
복사하는 게 아니고 각각 신경망?? 인데 (공용?)
State가 각각 다르기때문에 펼쳐놓고 보는 것(이해하기 쉽게?)
3개가 아님에 주의

6.2 Understanding recurrent neural networks

- ▶ RNN takes as input a sequence of vectors - 2D tensor of size (timesteps, input_features)
- ▶ set the **state** for the next step to be this previous output.
- ▶ *initial state* - **all-zero vector**

Listing 6.19 Pseudocode RNN

```
state_t = 0 # The state at t
for input_t in input_sequence:
    # Iterates over sequence elements
    output_t = f(input_t, state_t)
    state_t = output_t
    # The previous output becomes
    # the state for the next iteration.
```




6.2 Understanding recurrent neural networks

► f : input and state $\rightarrow W$ and U , and a bias vector

Listing 6.20 More detailed pseudocode for the RNN

```
state_t = 0 # h
for input_t in input_sequence:
    output_t = activation(dot(W, input_t) + dot(U, state_t) + b)
    state_t = output_t
```



6.2 Understanding recurrent neural networks

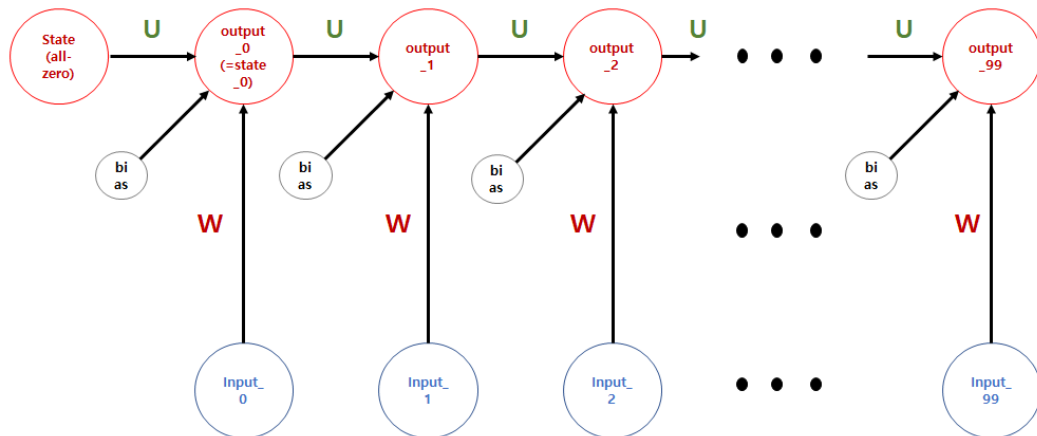
► naive Numpy implementation of the forward pass of the simple RNN.

Listing 6.21 Numpy implementation of a simple RNN

```
import numpy as np

timesteps = 100 # Number of timesteps in the input sequence
input_features = 32 # Dimensionality of the input feature space
output_features = 64 # Dimensionality of the output feature space
inputs = np.random.random((timesteps, input_features))
# Input data: random noise for the sake of the example
state_t = np.zeros((output_features,))
# Initial state: all-0 vector

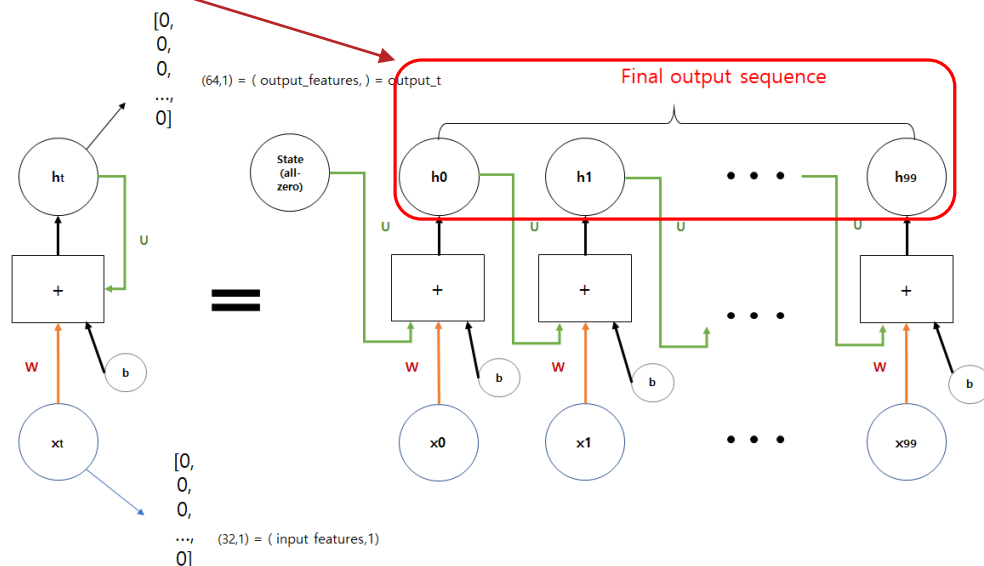

W = np.random.random((output_features, input_features)) # (64,32)
U = np.random.random((output_features, output_features)) # (64,64)
b = np.random.random((output_features,)) # (64,)
# 1 random weight matrices
```



6.2 Understanding recurrent neural networks

```
successive_outputs = []  
for input_t in inputs: # a vector of shape (input_features,)
    output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)  
    # Combines the input with the current state (the previous output)  
    successive_outputs.append(output_t) #Stores this output in a list  
    state_t = output_t  
    # Updates the state of the network for the next timestep  
final_output_sequence = np.concatenate(successive_outputs, axis=0)  
# The final output is a 2D tensor of  
# shape (timesteps, output_features), shape = (100, 64)
```

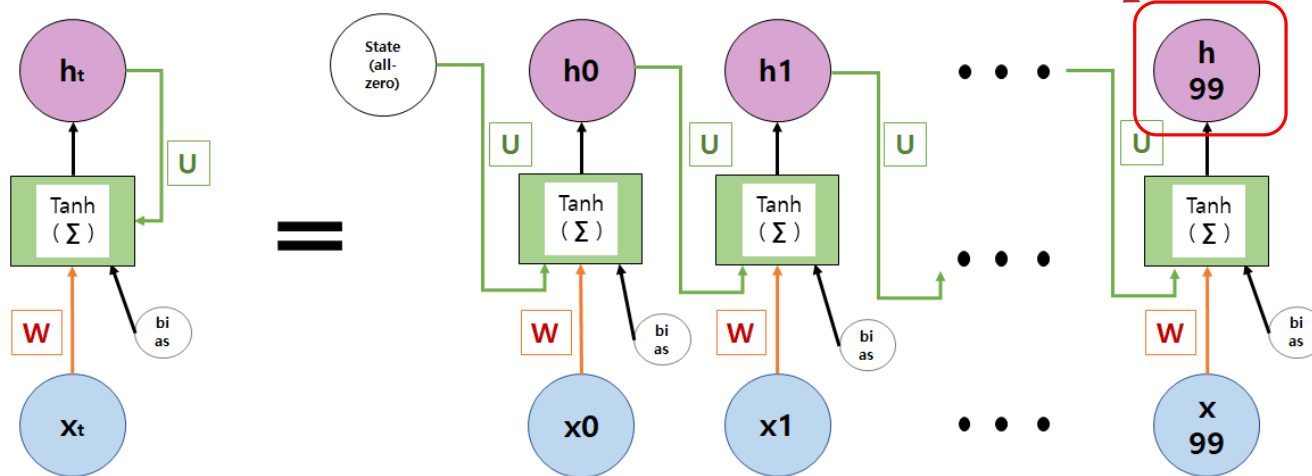
Setting `axis=0` concatenates along the row axis



6.2 Understanding recurrent neural networks

`output_t = np.tanh(np.dot(W, input_t) + np.dot(U, state_t) + b)`

NOTE The final output is a 2D tensor of shape (timesteps, output_features) at time t . Only the last output (`output_t`) at the end of the loop) is needed, because it already contains information about the entire sequence.



6.2 Understanding recurrent neural networks

6.2.1 A recurrent layer in Keras

- ▶ SimpleRNN layer:

```
from keras.layers import SimpleRNN
```

- ▶ There is one minor difference: SimpleRNN processes **batches of sequences**, like all other Keras layers

(timesteps, input_features) →

(batch_size, timesteps, input_features)

- ▶ two different modes of **return**

- ▶ (batch_size, timesteps, **output_features**) - the full sequences of successive outputs
- ▶ (batch_size, **output_features**) - only the last output for each input sequence

- ▶ These two modes are controlled by the **return_sequences** constructor argument.

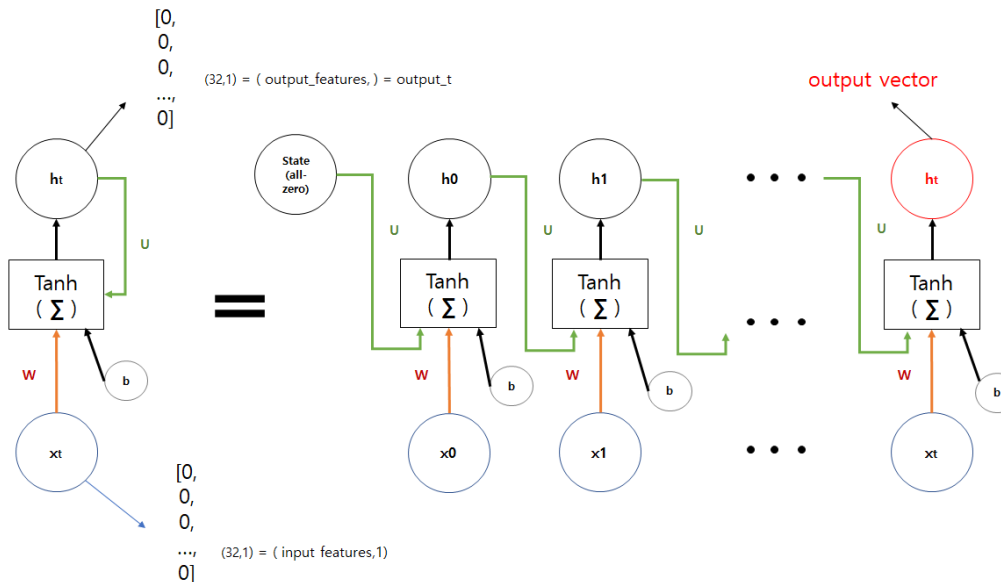
6.2 Understanding recurrent neural networks

- SimpleRNN and returns only the output at the **last timestep**:

```
>>> from keras.models import Sequential
>>> from keras.layers import Embedding, SimpleRNN
>>> model = Sequential()
>>> model.add(Embedding(10000, 32)) # (max_features, output dim)
>>> model.add(SimpleRNN(32))
>>> model.summary()
```

*10000은 가장 많이 사용되는 만 개
sparse는 10000개, 그걸 축약해서 32개로 embedding*

Layer (type)	Output Shape	Param #
embedding_22 (Embedding)	(None, None, 32)	320000
simplernn_10 (SimpleRNN)	(None, 32)	2080 = (32*32*2+32)

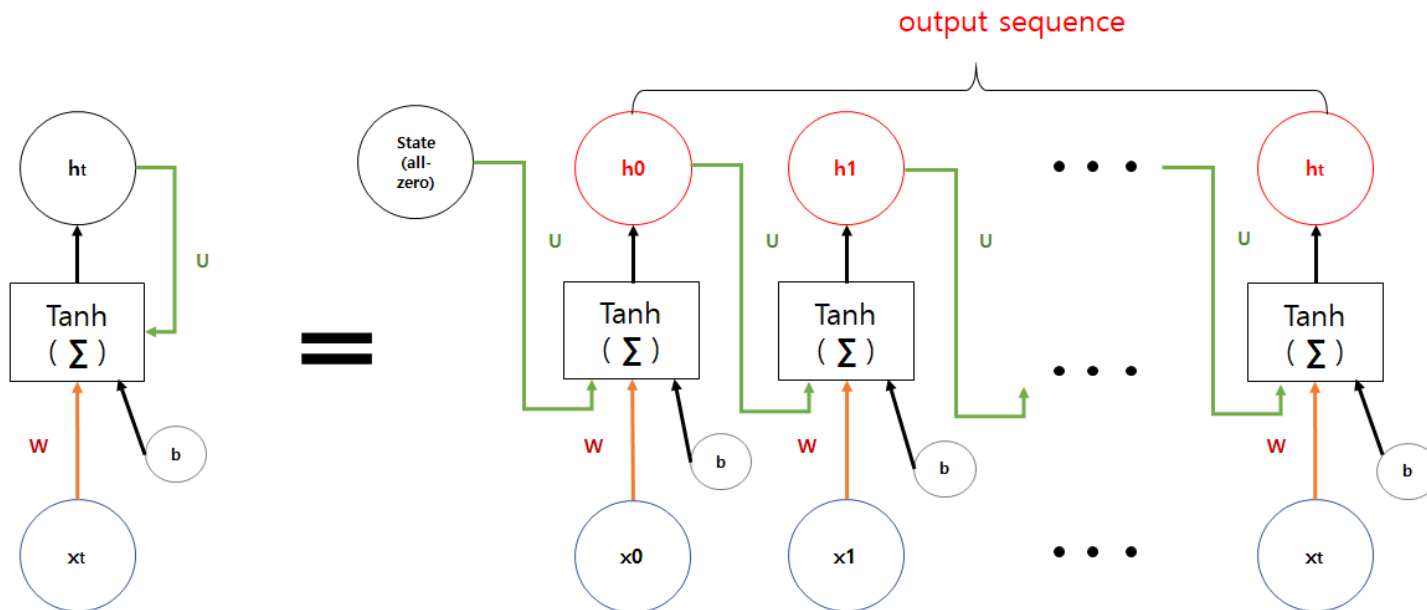


6.2 Understanding recurrent neural networks

- ▶ The following example returns the **full state sequence**:

```
>>> model = Sequential()  
>>> model.add(Embedding(10000, 32))  
>>> model.add(SimpleRNN(32, return_sequences=True)) 32개를 각각 다 저장한다?  
>>> model.summary()
```

Layer (type)	Output Shape	Param #
embedding_23 (Embedding)	(None, None, 32)	320000
simplernn_11 (SimpleRNN)	(None, None, 32)	2080




6.2 Understanding recurrent neural networks

- ▶ stack several recurrent layers:

```
>>> model = Sequential()
>>> model.add(Embedding(10000, 32))
>>> model.add(SimpleRNN(32, return_sequences=True))
>>> model.add(SimpleRNN(32, return_sequences=True))
>>> model.add(SimpleRNN(32, return_sequences=True))
>>> model.add(SimpleRNN(32)) # Last layer only returns the last output
>>> model.summary()
```

Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, None, 32)	320000
simple_rnn_3 (SimpleRNN)	(None, None, 32)	
simple_rnn_4 (SimpleRNN)	(None, None, 32)	2080
simple_rnn_5 (SimpleRNN)	(None, None, 32)	2080
simple_rnn_6 (SimpleRNN)	(None, 32)	2080



6.2 Understanding recurrent neural networks

- ▶ IMDB movie-review-classification problem - First, preprocess the data.

Listing 6.22 Preparing the IMDB data

```
from keras.datasets import imdb
from keras.preprocessing import sequence
max_features = 10000 # Number of words to consider as features
maxlen = 500 # Cuts off texts after this many words
batch_size = 32
(input_train, y_train), (input_test, y_test) =
    imdb.load_data(num_words=max_features)
print(len(input_train), 'train sequences') # 25000
print(len(input_test), 'test sequences') # 25000
input_train = sequence.pad_sequences(input_train, maxlen=maxlen)
input_test = sequence.pad_sequences(input_test, maxlen=maxlen)
print('input_train shape:', input_train.shape) # (25000, 500)
print('input_test shape:', input_test.shape) # (25000, 500)
```

6.2 *Understanding recurrent neural networks*

- ▶ Train an Embedding layer and a SimpleRNN layer.

Listing 6.23 Training the model with Embedding and SimpleRNN layers

```
from keras.layers import Dense
model = Sequential()
model.add(Embedding(max_features, 32))
model.add(SimpleRNN(32))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop', loss='binary_crossentropy',
              metrics=['acc'])
history = model.fit(input_train, y_train,
                    epochs=10, batch_size=128, validation_split=0.2) #20%
```



6.2 Understanding recurrent neural networks

Listing 6.24 Plotting results

```
import matplotlib.pyplot as plt
acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']
epochs = range(1, len(acc) + 1)

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()
plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()
plt.show()
```



6.2 Understanding recurrent neural networks

- ▶ In chapter 3, test accuracy - 88%
- ▶ recurrent network - 85% validation accuracy
- ▶ Inputs only the first 500 words - less information than the earlier baseline model.
- ▶ SimpleRNN - No good at processing long sequences, such as text (vanishing information).
- ▶ Other types of recurrent layers perform much better.

6.2 Understanding recurrent neural networks


6.2.2 A Understanding the LSTM and GRU layers

- ▶ LSTM and GRU - SimpleRNN has a major issue: **long-term dependencies** are impossible to learn.
- ▶ This is due to the *vanishing gradient problem*, an effect that is similar to what is observed with non-recurrent networks (feedforward networks) studied by Hochreiter, Schmidhuber, and Bengio in the early 1990s.
- ▶ Long Short-Term Memory (LSTM) algorithm was developed by Hochreiter and Schmidhuber in 1997.
- ▶ **Carry Track** (C) information across many timesteps to save information for later, thus preventing older signals from gradually vanishing during processing.


LSTM-GRU Architecture - overview

<https://towardsdatascience.com/illustrated-guide-to-lstms-and-gru-s-a-step-by-step-explanation-44e9eb85bf21>

LSTM



GRU



sigmoid



tanh



pointwise
multiplication

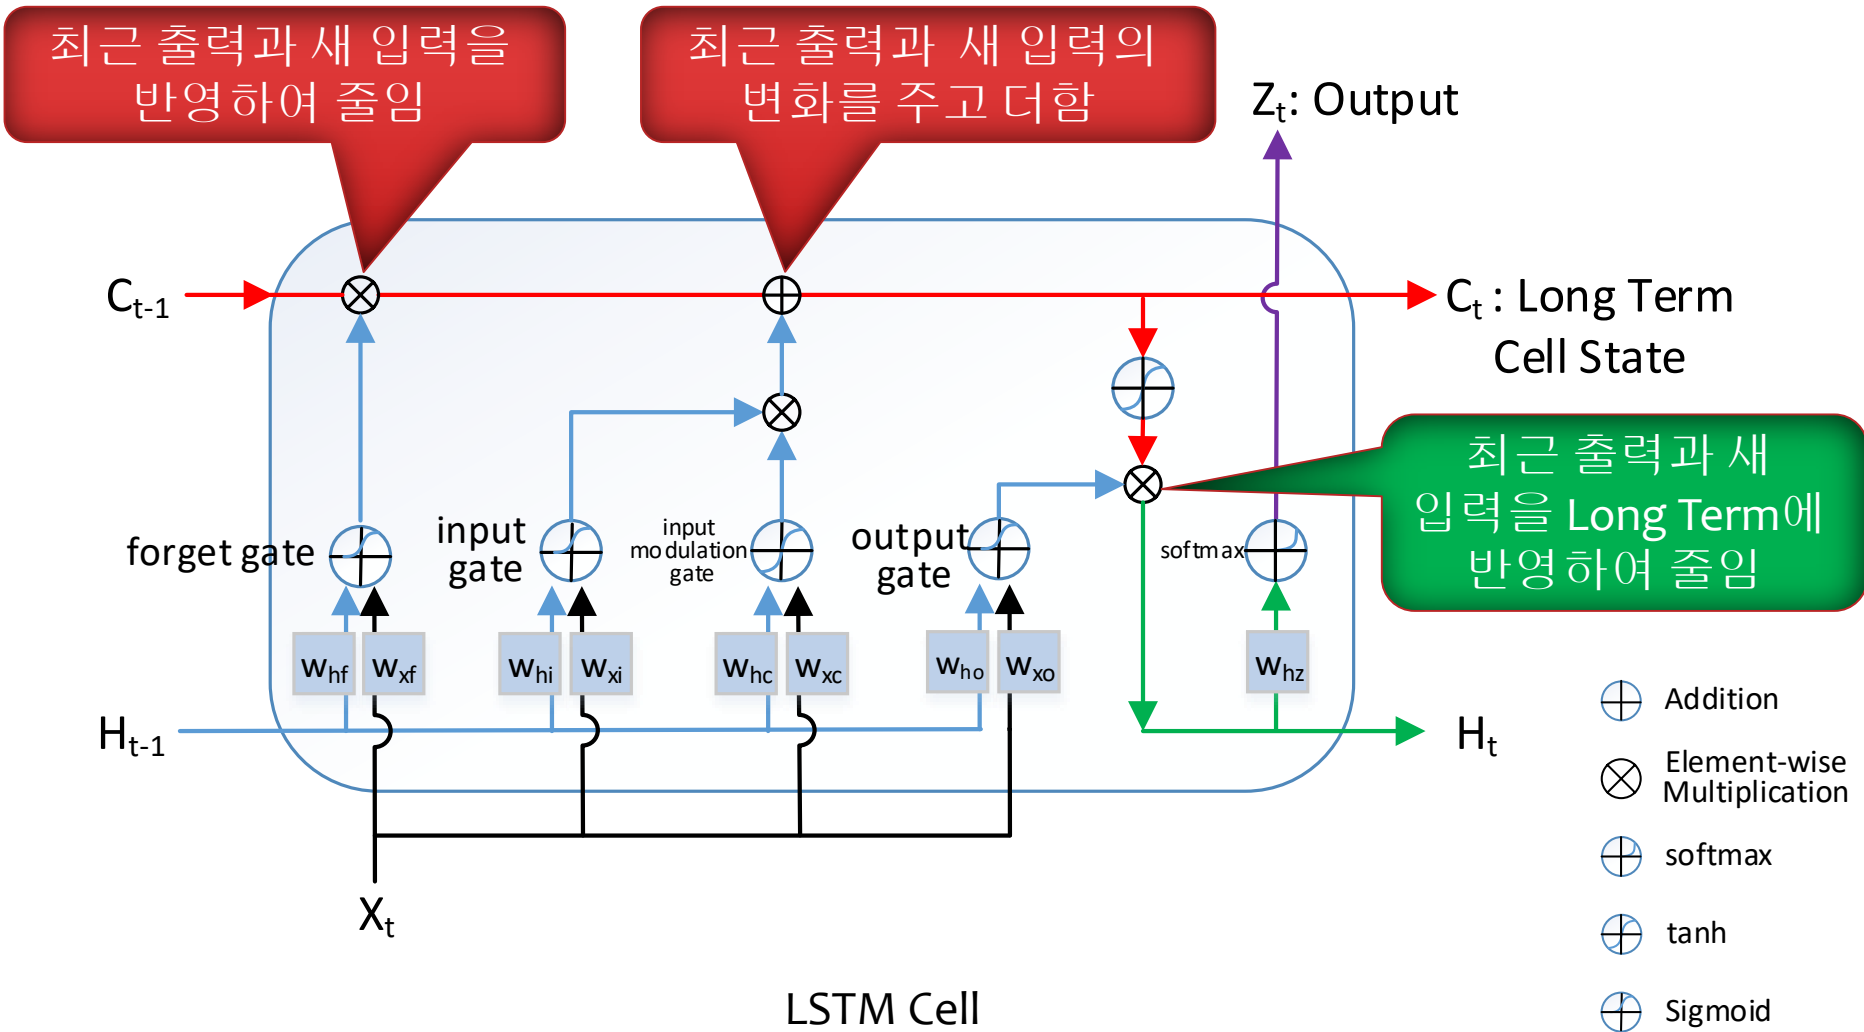


pointwise
addition




vector
concatenation

○ ○ ○ LSTM Architecture - overview ○ ○ ○




○ ○ ○ LSTM Architecture - overview ○ ○ ○




Forget gate operations

○ ○ ○ LSTM Architecture - overview ○ ○ ○




Input gate operations

○ ○ ○ LSTM Architecture - overview ○ ○ ○



Calculating cell state

○ ○ ○ LSTM Architecture - overview ○ ○ ○



output gate operations



6.2 Understanding recurrent neural networks

6.2.3 A concrete LSTM example in Keras

- ▶ set up a model using an LSTM layer and train it on the IMDB data (see figures 6.16 and 6.17).
- ▶ similar to the one with SimpleRNN - specify the output dimensionality of the LSTM layer; leave every other argument (there are many) at the Keras defaults.

Listing 6.27 Using the LSTM layer in Keras

```
from keras.layers import LSTM
model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='rmsprop',
              loss='binary_crossentropy', metrics=['acc'])
history = model.fit(input_train, y_train,
                    epochs=10, batch_size=128, validation_split=0.2)
```



6.2 Understanding recurrent neural networks

6.2.3A concrete LSTM example in Keras

- ▶ achieve up to **89%** validation accuracy with less vanishing-gradient problem—and slightly better than the fully connected approach from chapter 3
- ▶ **less data** than you were in chapter 3 by truncating sequences after 500 timesteps, whereas in chapter 3, you were considering full sequences.
- ▶ Why isn't LSTM performing better?
 - ▶ no effort to tune hyperparameters such as the **embeddings** dimensionality or the **LSTM output** dimensionality.
 - ▶ lack of **regularization**
 - ▶ analyzing the **global, long-term structure of the reviews** (what LSTM is good at) isn't helpful for a **sentiment-analysis** problem.
 - ▶ well solved by looking at what **words occur** in each review, and at what **frequency** in FCN
 - ▶ the strength of LSTM will become apparent: in particular, **question-answering** and **machine translation**

6.2.4 Wrapping up

- ▶ Now you understand the following:
 - What RNNs are and how they work
 - What LSTM is, and why it works better on long sequences than a naive RNN
 - How to use Keras RNN layers to process sequence data
- ▶ Next, advanced features of RNNs

6.3 Advanced use of recurrent neural networks

- ▶ **three advanced techniques** for improving the performance and generalization power of recurrent neural networks
- ▶ **temperature-forecasting** problem - timeseries of data points coming from sensors-temperature, air pressure, and humidity-to predict the temperature 24 hours after the last data point
- ▶ We'll cover the following techniques:
 - ▶ **Recurrent dropout** —to fight overfitting in recurrent layers
 - ▶ **Stacking** recurrent layers—This increases the representational power of the network (at the cost of higher computational loads).
 - ▶ **Bidirectional** recurrent layers—These present the same information to a recurrent network in different ways, increasing accuracy and mitigating forgetting issues.